



# CSC/CEG 3150 Tutorial

## Last Milestone in Assignment 3 & tips

XIAO Zigang

[zgxiao@cse.cuhk.edu.hk](mailto:zgxiao@cse.cuhk.edu.hk)

Department of Computer Science and Engineering  
The Chinese University of Hong Kong

November 18, 2008



# Outline

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

1 Outline

2 Milestone 3

3 Debug

4 Tips

5 Q & A

6 Acknowledgement



# Tasks in Milestone 3

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- **Objective:** Writing to an **existing** file
  - Truncation: **Discard** original data and write new content
  - Appendage: **Preserve** original data and write from the end of the target



# Truncation Mode work flow

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

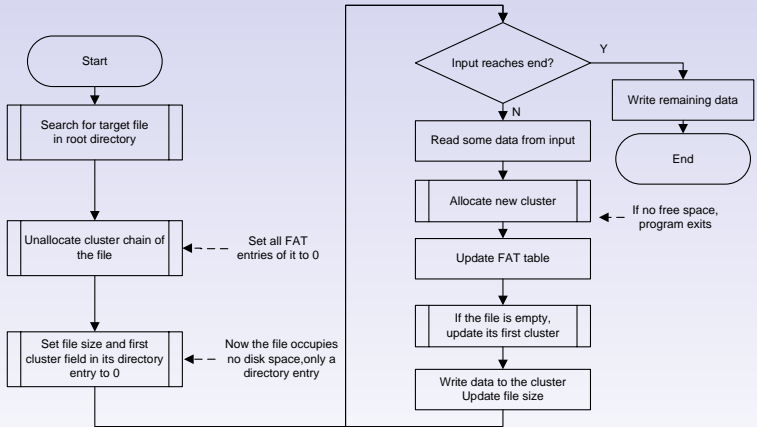


Figure: Illustration of truncation



- FSINFO data structure includes hints about where the OS can allocate new clusters.
- Its location is given in the boot sector

```
/* FSINFO, size=512B */  
#pragma pack(push,1)  
typedef struct FSInfo {  
    unsigned long FSI_LeadSig;           /* Signature (0x41615252) */  
    unsigned char FSI_Reserved1[480];    /* Not used */  
    unsigned long FSI_StrucSig;          /* Signature (0x61417272) */  
    unsigned long FSI_Free_Count;        /* Number of free clusters */  
    unsigned long FSI_Nxt_Free;          /* Next free cluster */  
    unsigned char FSI_Reserved2[12];     /* Not used */  
    unsigned long FSI_TrailSig;          /* Signature (0xAA550000) */  
}FSInfo;  
#pragma pack(pop)
```

Listing 1: FSINFO structure



# Cluster allocation

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- Different allocation scheme in implementations
- You should implement a **circular, next-available** scheme
- The next free cluster information is stored in FSINFO structure
  - i.e. *FSI\_Nxt\_Free*
  - Note: it stores the cluster value **just assigned**
- Remember to update FSINFO after a cluster is allocated!



Figure: Illustration of truncation



# Appendage Mode work flow

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

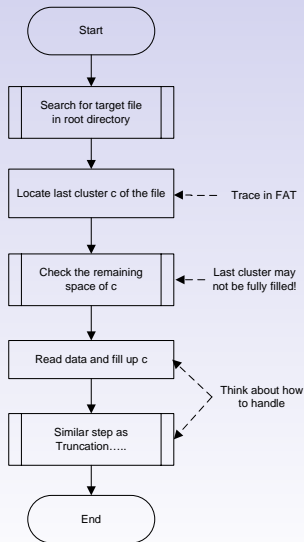


Figure: Illustration of appendage



# How to debug your program?

CSC/CEG  
3150  
Tutorial

- Here is one scenario you may follow:

```
# create a disk image file, size = 64M
$ dd if=/dev/zero of=./fs bs=1M count=64
$ format as FAT32 filesystem(you may also use /dev/ram[x] here)
$ /sbin/mkfs.vfat -v -F 32 -f 2 -S 512 -s 1 -R 32 ./fs
# you may backup this file here for later use

# mount to a certain mount point
$ mkdir /mnt/rd
$ mount -t vfat -o loop,umask=000 ./fs /mnt/rd
# You can write a script/make rule to do stuff listed above

# make some change here...
$ echo "void main(){}" > /mnt/rd/TEST.C
$ touch /mnt/rd/MAKEFILE
# synchronize change
$ sync

# check your program
$ make
$ ./a3 -m 1 -d ./fs
$ sync
# ... (check your output)
```

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement





# utilize make

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- **make** is very convenient in this assignment, e.g.
  - include **DEBUG** switch in your program, i.e.

```
#ifdef DEBUG
    printf("program reaches here...\n");
    ...
#endif
```

- when compile your program with  
**#define DEBUG**  
or  
**gcc <other-options> -DDEBUG**  
The codes between **#ifdef/#endif** will be included.
- You can add a **debug rule** in your makefile, then you may use:  
**\$ make debug**  
to compile a debug version



# Example of test batch

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- You are strongly encouraged to write script/make rule to initialize your disk image
- Please be reminded that `mkfs.vfat` will NOT help you to clear original content in the disk image

```
$ cat ./initfs
#!/bin/bash
dd if=/dev/zero of=filesystem bs=1M count=64
/sbin/mkfs.vfat -v -F 32 -f 2 -S 512 -s 1 -R 32 ./filesystem
mount -t vfat ./filesystem ./rd -o loop,umask=000 -v

echo 123 > ./rd/A.TXT
# file "ascii" contains some ascii characters
cat ascii > ./rd/B.TXT
dd if=/dev/urandom of=./rd/C.TXT bs=64 count=1
# file "raw" contains some raw contents
cat raw > ./rd/D.TXT

sync
```

Listing 2: Example of script



# Sample output

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

```
$ ./a3 -m 1 -d ./filesystem
FILENAME          FILE SIZE          STARTING CLUSTER #
-----
A.TXT             4                  3
B.TXT             800                 4
C.TXT             100                 6
D.TXT             800                 7

Total number of entries = 4
$ ./a3 -m 2 -d ./filesystem -t C.TXT | xxd
0000000: cd5e 10e7 e7ec 1988 006b 6974 7e5b f3a5  .^.....kit~[..
0000010: 78cd 7ab3 daa9 2687 8814 269d b2c3 fd7f  x.z...&...&....
...
$ cat rd/A.TXT
123
$ echo "hello :)" > input_file.txt
# use truncation mode here
$ umount filesystem
$ ./a3 -m 3 -d filesystem -t A.TXT -w t -i input_file.txt && sync
$ mount -t vfat ./filesystem ./rd -o loop,umask=000 -v
$ cat rd/A.TXT
hello :)
# use appendage mode here
$ umount filesystem
$ ./a3 -m 3 -d filesystem -t A.TXT -w a -i input_file.txt && sync
$ mount -t vfat ./filesystem ./rd -o loop,umask=000 -v
$ ./a3 -m 2 -d filesystem -t A.TXT
hello :)
hello :)
```

Listing 3: Sample output



# Some facts

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- Where is the root directory?
  - Check BootEntry structure
- How many Directory Entries in a cluster?
  - $= \text{cluster\_size} / \text{DirEntry\_size}$
- What is the usage of the reserved sectors, other than BootEntry,FSINFO?
  - Reserved for system use(e.g. backup, upgrade)
- What is a DirEntry?
  - It contains the name and metadata for a file or directory. Each file or directory is allocated one DirEntry, it is located in the clusters allocated to the file's parent directory



# Some facts(cont.)

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- How do we know if a DirEntry is available?
  - Allocation status of a directory entry is determined by using the first byte:
    - 0x00 or 0x0e : unallocated
    - Others: allocated
- What is the size of a directory?
  - The size field in the DirEntry of a directory is not used and should always be 0.
  - The only way to determine the size of the directory is to follow the cluster chain in FAT
- What is the usage of entry 0 and 1 in FAT table?
  - Cluster starts from index 2
  - Generally, 0 records media type, 1 records dirty status of file system



# Hints

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- ① You are not required to update *create time*, *access time*, *etc.* fields in Directory entry in this assignment.
- ② After changes are made to the file system(e.g. create/remove a file), you may find that the file system read by your program is not updated.
- ③ Kernel keeps data in memory in order to avoid slow disk I/O
  - Data will be written back “gradually”
- ④ Use `sync` to flush the data
- ⑤ or use `sync()` in your program
- ⑥ If `sync` does not work, try `umount` and then `mount` again.



# Q & A Session

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

# Thank You

Now you should understand FAT32  
- Have fun in assignment 3 -



# Acknowledgement

CSC/CEG  
3150  
Tutorial

Outline

Milestone 3

Debug

Tips

Q & A

Acknowledgement

- Some materials and pictures are from last year's tutorial notes made by *Mr. Cheong Chi Hong*
- The style of this slide is adapted from the template made by *HUANG Zheng-hua* in *Wuhan University*
- Google “vfs” to learn more about filesystem in Linux